

```
In [136... import pandas as pd
import numpy as np
```

```
In [137... data=pd.read_csv('team.csv')
```

```
In [138... from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
dfle=data
dfle.TEAM=le.fit_transform(dfle.TEAM)
dfle
```

```
Out[138]:
```

| | TEAM | YEAR |
|---|------|------|
| 0 | 0 | 2000 |
| 1 | 1 | 2002 |
| 2 | 2 | 2003 |
| 3 | 3 | 2004 |
| 4 | 0 | 2005 |
| 5 | 2 | 2006 |
| 6 | 1 | 2007 |
| 7 | 0 | 2008 |
| 8 | 3 | 2009 |

```
In [139... from sklearn.preprocessing import OneHotEncoder
```

```
In [140... enc=OneHotEncoder()
enc_data=pd.DataFrame(enc.fit_transform(dfle[['TEAM']]).toarray())
enc_data
```

```
Out[140]:
```

| | 0 | 1 | 2 | 3 |
|---|-----|-----|-----|-----|
| 0 | 1.0 | 0.0 | 0.0 | 0.0 |
| 1 | 0.0 | 1.0 | 0.0 | 0.0 |
| 2 | 0.0 | 0.0 | 1.0 | 0.0 |
| 3 | 0.0 | 0.0 | 0.0 | 1.0 |
| 4 | 1.0 | 0.0 | 0.0 | 0.0 |
| 5 | 0.0 | 0.0 | 1.0 | 0.0 |
| 6 | 0.0 | 1.0 | 0.0 | 0.0 |
| 7 | 1.0 | 0.0 | 0.0 | 0.0 |
| 8 | 0.0 | 0.0 | 0.0 | 1.0 |

```
In [142... abc=dfle.join(enc_data)
```

```
In [143... final=abc.drop(['TEAM'],axis='columns')
final
```

```
Out[143]:
```

| | YEAR | 0 | 1 | 2 | 3 |
|---|------|-----|-----|-----|-----|
| 0 | 2000 | 1.0 | 0.0 | 0.0 | 0.0 |
| 1 | 2002 | 0.0 | 1.0 | 0.0 | 0.0 |
| 2 | 2003 | 0.0 | 0.0 | 1.0 | 0.0 |
| 3 | 2004 | 0.0 | 0.0 | 0.0 | 1.0 |
| 4 | 2005 | 1.0 | 0.0 | 0.0 | 0.0 |
| 5 | 2006 | 0.0 | 0.0 | 1.0 | 0.0 |
| 6 | 2007 | 0.0 | 1.0 | 0.0 | 0.0 |
| 7 | 2008 | 1.0 | 0.0 | 0.0 | 0.0 |
| 8 | 2009 | 0.0 | 0.0 | 0.0 | 1.0 |

In []:

```
In [1]: import pandas as pd  
import numpy as np
```

```
In [2]: data=pd.read_csv('abcde.csv')
```

```
In [4]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(data[['age']],data.results,test_size=0.5,random_state=10)
```

```
In [5]: from sklearn.linear_model import LogisticRegression  
model=LogisticRegression()  
model.fit(x_train,y_train)
```

```
Out[5]: LogisticRegression()  
In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.
```

```
In [6]: y_pred=model.predict(x_test)
```

```
In [7]: model.score(x_test,y_test)
```

```
Out[7]: 0.9285714285714286
```

```
In [ ]:
```

In [1]: Write Python program to implement classifier using support vector machines.
 1) Installation packages , Loading dataset, Show data frame , Create confusion_matrix
 2) Describe accuracy_score, precision_score, recall_score, f1_score using confusion matrix

In [2]: `from sklearn.datasets import load_iris`
`iris=load_iris()`

In [7]: `df=pd.DataFrame(iris.data,columns=iris.feature_names)`

In [8]: `df['target']=iris.target`
`df.head()`

Out[8]:

| | sepal length (cm) | sepal width (cm) | petal length (cm) | petal width (cm) | target |
|---|-------------------|------------------|-------------------|------------------|--------|
| 0 | 5.1 | 3.5 | 1.4 | 0.2 | 0 |
| 1 | 4.9 | 3.0 | 1.4 | 0.2 | 0 |
| 2 | 4.7 | 3.2 | 1.3 | 0.2 | 0 |
| 3 | 4.6 | 3.1 | 1.5 | 0.2 | 0 |
| 4 | 5.0 | 3.6 | 1.4 | 0.2 | 0 |

In [11]: `x=df.drop(['target'],axis='columns')`
`y=df.target`

In [12]: `from sklearn.model_selection import train_test_split`
`x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=0)`

In [13]: `from sklearn.svm import SVC`
`clr=SVC()`
`clr.fit(x_train,y_train)`

Out[13]: SVC()
 In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
 On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [14]: `clr.score(x_test,y_test)`

Out[14]: 1.0

In [30]: `y_pred=clr.predict(x_test)`
`y_pred`

Out[30]: array([2, 1, 0, 2, 0, 2, 0, 1, 1, 1, 2, 1, 1, 1, 1, 0, 1, 1, 0, 0, 2, 1,
 0, 0, 2, 0, 0, 1, 1, 0])

In [31]: `from sklearn.metrics import classification_report`
`classification_report(y_test,y_pred)`

Out[31]:

| | precision | recall | f1-score | support | | 0 | 1.00 | 1.00 | 1.00 |
|-------|--------------|--------|----------|---------|------|-----------|------|------|------|
| 11\n | 1 | 1.00 | 1.00 | 1.00 | 13\n | 2 | 1.00 | 1.00 | 1.00 |
| 6\n\n | accuracy | | | 1.00 | 30\n | macro avg | 1.00 | 1.00 | 1.00 |
| 30\n | weighted avg | 1.00 | 1.00 | 1.00 | 30\n | | | | |

In []:

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
In [3]: from sklearn.datasets import load_iris
iris=load_iris()
```

```
In [4]: df=pd.DataFrame(iris.data,columns=iris.feature_names)
df['target']=iris.target
```

```
In [5]: x=df.drop(['target'],axis=1)
y=df['target']
```

```
In [6]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=2)
```

```
In [7]: from sklearn.svm import SVC
clr=SVC()
clr.fit(x_train,y_train)
```

Out[7]: SVC()
**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [8]: y_train_pred=clr.predict(x_train)
y_test_pred=clr.predict(x_test)
```

```
In [15]: from sklearn.metrics import confusion_matrix
train_cm=confusion_matrix(y_train,y_train_pred)
train_cm
```

Out[15]: array([[36, 0, 0],
 [0, 41, 1],
 [0, 2, 40]], dtype=int64)

```
In [17]: test_cm=confusion_matrix(y_test,y_test_pred)
test_cm
```


Out[17]: array([[14, 0, 0],
 [0, 7, 1],
 [0, 0, 8]], dtype=int64)

```
In [31]: #Visualize_Confusion_matrix
import seaborn as sns

plt.figure(figsize=(10,4))
sns.heatmap(train_cm,annot=True,fmt='d',cmap='Blues',
            xticklabels=iris.target_names,
            yticklabels=iris.target_names)

plt.title('Confusion Matrix - Test Set')
plt.xlabel('Predicted')
plt.ylabel('Actual')

plt.tight_layout()
plt.show()
```



```
In [33]: from sklearn.metrics import classification_report
classification_report(y_test,y_test_pred)
```

```
Out[33]: '          precision    recall  f1-score   support\n\n 14\n 1          1.00      0.88      0.93         8\n 8\n 8\n 30\n 6\n 30\nweighted avg          0.97      0.97      0.97\n\n              0          1          1          1\n              2          0.89          1.00          0.94\nmacro avg          0.96          0.96          0.9
```

In []:

```
In [2]: import pandas as pd
import numpy as np
```

```
In [3]: from sklearn.datasets import load_iris
data=load_iris()
```

```
In [4]: data.feature_names
```

```
Out[4]: ['sepal length (cm)',
'sepal width (cm)',
'petal length (cm)',
'petal width (cm)']
```

```
In [5]: data.target_names
```

```
Out[5]: array(['setosa', 'versicolor', 'virginica'], dtype='<U10')
```

```
In [7]: x=data.data
y=data.target
```

```
In [8]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=10)
```

```
In [9]: from sklearn.tree import DecisionTreeClassifier
sc=DecisionTreeClassifier(criterion='entropy')
sc.fit(x_train,y_train)
```

```
Out[9]: DecisionTreeClassifier(criterion='entropy')
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [10]: y_pred=sc.predict(x_test)
```

```
In [16]: from sklearn.metrics import confusion_matrix,accuracy_score,classification_report
confusion_matrix(y_test,y_pred)
```

```
Out[16]: array([[10,  0,  0],
[ 0, 11,  2],
[ 0,  0,  7]], dtype=int64)
```

```
In [17]: accuracy_score(y_test,y_pred)
```

```
Out[17]: 0.9333333333333333
```

```
In [18]: classification_report(y_test,y_pred)
```

```
Out[18]: '
          precision    recall  f1-score   support\n\n
 0          1.00      1.00      1.00         10\n
 1          0.78      1.00      0.88         13\n
 2          0.93      0.93      0.93         30\n
-----
 overall accuracy      0.93      0.93      0.93        53\n
weighted avg          0.95      0.93      0.93        53'
```

```
In [ ]:
```

```
In [1]: import numpy as np
import pandas as pd
```

```
In [2]: from sklearn.datasets import load_breast_cancer
data=load_breast_cancer()
```

```
In [3]: data.target_names
```

```
Out[3]: array(['malignant', 'benign'], dtype='<U9')
```

```
In [6]: df=pd.DataFrame(np.c_[data.data,data.target],columns=[list(data.feature_names)+'target'])
```

```
In [8]: df.shape
```

```
Out[8]: (569, 31)
```

```
In [10]: x=df.iloc[:,0:-1]
y=df.iloc[:, -1]
```

```
In [13]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=1)
```

```
In [14]: from sklearn.preprocessing import StandardScaler
sc=StandardScaler()
xtrain=sc.fit_transform(x_train)
xtest=sc.transform(x_test)
```

```
In [15]: from sklearn.naive_bayes import GaussianNB
clr=GaussianNB()
clr.fit(xtrain,y_train)
```

```
Out[15]: GaussianNB()
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [16]: y_pred=clr.predict(xtest)
```

```
In [17]: y_pred
```

```
Out[17]: array([0., 0., 1., 0., 1., 0., 0., 0., 1., 1., 1., 0., 0., 1., 1., 1., 1.,
        1., 1., 0., 1., 1., 0., 1., 0., 1., 1., 0., 0., 0., 1., 0., 0.,
        1., 1., 0., 1., 1., 1., 1., 1., 1., 1., 1., 0., 1., 1., 1., 0., 0.,
        0., 1., 1., 1., 1., 1., 0., 1., 1., 1., 0., 1., 1., 1., 1., 0.,
        1., 1., 1., 1., 1., 0., 1., 0., 0., 1., 1., 0., 1., 0., 1.,
        1., 0., 1., 0., 1., 1., 0., 1., 1., 0., 0., 1., 1., 1., 1., 1., 1.,
        1., 1., 1., 1., 1., 1., 0., 0., 0., 0., 1., 1.])
```

```
In [19]: from sklearn.metrics import accuracy_score
print('Accuracy :',accuracy_score(y_test,y_pred))
```

```
Accuracy : 0.9473684210526315
```

```
In [ ]:
```



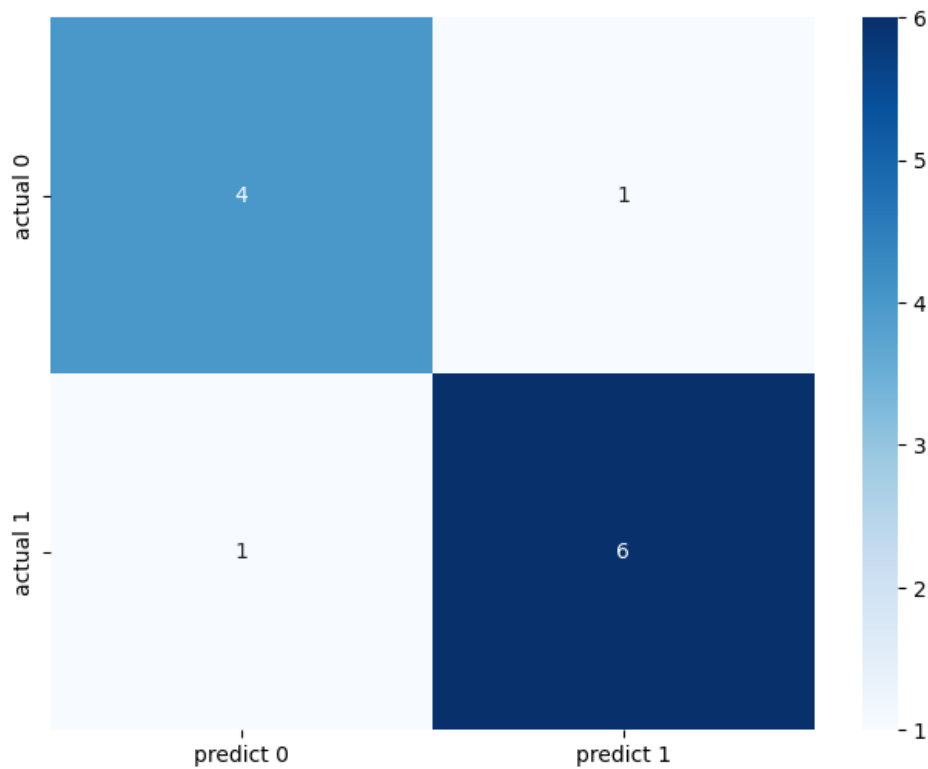
```
In [1]: from sklearn.metrics import confusion_matrix, classification_report
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: actual=[1,0,1,0,1,0,1,0,1,1,0]
predict=[1,0,1,0,1,0,1,1,1,1,0,0]
```

```
In [4]: cm=confusion_matrix(actual,predict)
```

```
In [10]: plt.figure(figsize=(8,6))
sns.heatmap(cm,annot=True,fmt='d',cmap='Blues',
            xticklabels=['predict 0','predict 1'],
            yticklabels=['actual 0','actual 1'])
```

Out[10]: <Axes: >



```
In [12]: classification_report(actual,predict)
```

```
Out[12]: '
          precision    recall  f1-score   support\n\n
 0          0.80      0.80      0.80         7\n\n
 1          0.83      0.83      0.83        12\n\n
 macro avg   0.83      0.83      0.83        19\n\n
 weighted avg 0.83      0.83      0.83        19'
```

```
In [ ]:
```

```
In [1]: #Random_Forest
import pandas as pd
import numpy as np
from sklearn import metrics
```

```
In [2]: data=pd.read_csv('user_data.csv')
```

```
In [3]: x=data.iloc[:,[2,3]].values
y=data.iloc[:,4].values
```

```
In [4]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=2)
```

```
In [5]: from sklearn.preprocessing import StandardScaler
sc=StandardScaler()
x_train=sc.fit_transform(x_train)
x_test=sc.transform(x_test)
```

```
In [6]: data
```

```
Out[6]:
```

| | User ID | Gender | Age | EstimatedSalary | Purchased |
|-----|----------|--------|-----|-----------------|-----------|
| 0 | 15624510 | Male | 19 | 19000 | 0 |
| 1 | 15810944 | Male | 35 | 20000 | 0 |
| 2 | 15668575 | Female | 26 | 43000 | 0 |
| 3 | 15603246 | Female | 27 | 57000 | 0 |
| 4 | 15804002 | Male | 19 | 76000 | 0 |
| ... | ... | ... | ... | ... | ... |
| 395 | 15691863 | Female | 46 | 41000 | 1 |
| 396 | 15706071 | Male | 51 | 23000 | 1 |
| 397 | 15654296 | Female | 50 | 20000 | 1 |
| 398 | 15755018 | Male | 36 | 33000 | 0 |
| 399 | 15594041 | Female | 49 | 36000 | 1 |

400 rows × 5 columns

```
In [7]: from sklearn.ensemble import RandomForestClassifier
clr=RandomForestClassifier(n_estimators=10,criterion='entropy')
clr.fit(x_train,y_train)
```

```
Out[7]: RandomForestClassifier(criterion='entropy', n_estimators=10)

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.
```

```
In [9]: y_pred=clr.predict(x_test)
y_pred
```

```
Out[9]: array([0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 1, 1, 0, 0, 0,
               0, 0, 0, 1, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1,
               0, 1, 0, 1, 1, 0, 0, 1, 0, 1, 1, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0, 1,
               1, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0], dtype=int64)
```

```
In [12]: from sklearn.metrics import confusion_matrix,classification_report
cm=confusion_matrix(y_test,y_pred)
cm
```

```
Out[12]: array([[43,  5],
               [ 6, 26]], dtype=int64)
```

```
In [13]: from sklearn.metrics import classification_report
classification_report(y_test,y_pred)
```

```
Out[13]: '              precision    recall  f1-score   support\n\n 48\n      1          0.84          0.81          0.83         32\n\n accuracy          0.88          0.90          0.89\n 86      80\n macro avg          0.86          0.85          0.86\n 0.86      80\nweighted avg          0.86          0.86          0.86'
```

```
In [ ]: #Visualization Not Possible The Program Going To Confusion
```

```
In [1]: import pandas as pd
from sklearn.preprocessing import MinMaxScaler
from matplotlib import pyplot as plt
```


```
In [2]: df=pd.read_csv('Book1.csv')
df.head()
```

```
Out[2]:
```

| | name | rollno | marks |
|---|------|--------|-------|
| 0 | A | 40 | 65 |
| 1 | B | 41 | 63 |
| 2 | C | 43 | 64 |
| 3 | D | 39 | 80 |
| 4 | E | 36 | 156 |

```
In [3]: plt.scatter(df.rollno,df['marks'])
plt.xlabel('rollno')
plt.ylabel('marks')
```

```
Out[3]: Text(0, 0.5, 'marks')
```



```
In [5]: from sklearn.cluster import KMeans
km=KMeans(n_clusters=3)
pred=km.fit_predict(df[['rollno','marks']])
pred
```


C:\Users\Admin\anaconda3\lib\site-packages\sklearn\cluster_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning

C:\Users\Admin\anaconda3\lib\site-packages\sklearn\cluster_kmeans.py:1382: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=1.

```
Out[5]: array([2, 2, 2, 2, 1, 1, 1, 0, 0, 0, 0, 0, 2, 2, 0, 0, 1, 1, 1, 1])
```

```
In [7]: df['cluster']=pred
df.head()
df1=df[df.cluster==0]
df2=df[df.cluster==1]
df3=df[df.cluster==2]
plt.scatter(df1.rollno,df1['marks'],color='green')
plt.scatter(df2.rollno,df2['marks'],color='red')
plt.scatter(df3.rollno,df3['marks'],color='blue')
plt.xlabel('rollno')
plt.ylabel('marks')
```

Out[7]: Text(0, 0.5, 'marks')



```
In [9]: scale=MinMaxScaler()
scale.fit(df[['marks']])
df['marks']=scale.transform(df[['marks']])

scale.fit(df[['rollno']])
df['rollno']=scale.transform(df[['rollno']])
```

```
In [13]: df=df.drop(['cluster'],axis='columns')
df['cluster']=pred
df.head()
```

Out[13]:

| | name | rollno | marks | cluster |
|---|------|----------|----------|---------|
| 0 | A | 0.823529 | 0.170940 | 2 |
| 1 | B | 0.882353 | 0.153846 | 2 |
| 2 | C | 1.000000 | 0.162393 | 2 |
| 3 | D | 0.764706 | 0.299145 | 2 |
| 4 | E | 0.588235 | 0.948718 | 1 |


```
In [14]: km.cluster_centers_
```

Out[14]: array([[28.42857143, 52.42857143],
[38.28571429, 150.],
[36.5 , 72.]])

```
In [15]: plt.scatter(df1.rollno,df1['marks'],color='green')
plt.scatter(df2.rollno,df2['marks'],color='red')
plt.scatter(df3.rollno,df3['marks'],color='blue')

plt.scatter(km.cluster_centers_[0,0],km.cluster_centers_[0,1],color='black',marker='*')
plt.xlabel('rollno')
plt.ylabel('marks')
```

Out[15]: Text(0, 0.5, 'marks')



In []:

```
In [1]: #KNN
import pandas as pd
import numpy as np
```

```
In [2]: data=pd.read_csv('User_Data.csv')
```

```
In [4]: x=data.iloc[:,[2,3]].values
y=data.iloc[:,4].values
```

```
In [5]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=2)
```

```
In [6]: from sklearn.preprocessing import StandardScaler
sc=StandardScaler()
xtrain=sc.fit_transform(x_train)
xtest=sc.transform(x_test)
```

```
In [8]: from sklearn.neighbors import KNeighborsClassifier
clr=KNeighborsClassifier(n_neighbors=5)
clr.fit(xtrain,y_train)
```

```
Out[8]: KNeighborsClassifier()
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [9]: y_pred=clr.predict(xtest)
y_pred
```

```
Out[9]: array([0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0,
               0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0,
               0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0, 1,
               1, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 1, 0], dtype=int64)
```

```
In [10]: from sklearn.metrics import accuracy_score,confusion_matrix,classification_report
cm=confusion_matrix(y_test,y_pred)
cm
```

```
Out[10]: array([[46,  2],
                [ 7, 25]], dtype=int64)
```

```
In [11]: classification_report(y_test,y_pred)
```

```
Out[11]: '
           precision    recall  f1-score   support\n\n
  0               0.87      0.96      0.91         32\n\n
  1               0.93      0.78      0.85         48\n\n
 accuracy                   0.89\n
 macro avg              0.90      0.87      0.88         80\n
weighted avg              0.89      0.89      0.89         80\n'
```

```
In [12]: accuracy_score(y_test,y_pred)
```

```
Out[12]: 0.8875
```

```
In [ ]: #Visualization_Not_Possible
```